

K means clustering

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

In [2]:

```
df=pd.read_csv(r"C:\Users\jyothi reddy\Downloads\Income.csv")
df
```

Out[2]:

	Gender	Age	Income(\$)
0	Male	19	15
1	Male	21	15
2	Female	20	16
3	Female	23	16
4	Female	31	17
...
195	Female	35	120
196	Female	45	126
197	Male	32	126
198	Male	32	137
199	Male	30	137

200 rows × 3 columns

In [3]:

```
df.head()
```

Out[3]:


	Gender	Age	Income(\$)
0	Male	19	15
1	Male	21	15
2	Female	20	16
3	Female	23	16
4	Female	31	17

In [4]:

```
plt.scatter(df["Age"],df["Income($)"])  
plt.xlabel("Age")  
plt.ylabel("Income($)")
```

Out[4]:

Text(0, 0.5, 'Income(\$)')



In [5]:

```
from sklearn.cluster import KMeans  
km=KMeans()  
km
```

Out[5]:

▼ KMeans
KMeans()

In [6]:



```
y_predicted=km.fit_predict(df[["Age", "Income($)"]])
y_predicted
```

C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

```
warnings.warn(
```

Out[6]:

```
array([2, 2, 2, 2, 2, 2, 2, 2, 7, 2, 7, 2, 7, 2, 2, 2, 2, 2, 7, 2, 2, 2,
       7, 2, 7, 2, 7, 2, 2, 2, 7, 2, 7, 2, 7, 2, 7, 2, 2, 2, 7, 2, 7, 2,
       7, 2, 7, 2, 2, 2, 7, 2, 2, 7, 7, 7, 7, 3, 0, 7, 3, 0, 3, 7, 3, 0,
       6, 3, 0, 0, 3, 6, 3, 3, 3, 0, 6, 6, 0, 6, 3, 6, 3, 6, 0, 6, 6, 0,
       0, 6, 3, 0, 6, 6, 0, 0, 6, 0, 6, 0, 0, 6, 3, 0, 6, 0, 3, 6, 3, 3,
       3, 0, 6, 0, 0, 0, 3, 6, 6, 6, 0, 6, 6, 6, 0, 1, 6, 1, 6, 1, 6, 1,
       1, 1, 1, 1, 6, 1, 1, 1, 6, 1, 1, 1, 1, 1, 6, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 6, 1, 1, 1, 5, 1, 5, 1, 1, 1, 5, 1, 1, 1, 5, 1,
       5, 1, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 4, 4, 4, 4, 4, 4,
       4, 4])
```

In [7]:



```
df["cluster"]=y_predicted
df.head()
```

Out[7]:


	Gender	Age	Income(\$)	cluster
0	Male	19	15	2
1	Male	21	15	2
2	Female	20	16	2
3	Female	23	16	2
4	Female	31	17	2

In [8]:

```
df1=df[df.cluster==0]
df2=df[df.cluster==1]
df3=df[df.cluster==2]
plt.scatter(df1["Age"],df1["Income($)"],color="red")
plt.scatter(df2["Age"],df2["Income($)"],color="green")
plt.scatter(df3["Age"],df3["Income($)"],color="blue")
plt.xlabel("Age")
plt.ylabel("Income($)")
```

Out[8]:

Text(0, 0.5, 'Income(\$)')



In [9]:



```
from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()
scaler.fit(df[["Income($)"]])
df["Income($)"]=scaler.transform(df[["Income($)"]])
df.head()
```

Out[9]:

	Gender	Age	Income(\$)	cluster
0	Male	19	0.000000	2
1	Male	21	0.000000	2
2	Female	20	0.008197	2
3	Female	23	0.008197	2
4	Female	31	0.016393	2

In [10]:



```
scaler.fit(df[["Age"]])
df["Age"]=scaler.transform(df[["Age"]])
df.head()
```

Out[10]:

	Gender	Age	Income(\$)	cluster
0	Male	0.019231	0.000000	2
1	Male	0.057692	0.000000	2
2	Female	0.038462	0.008197	2
3	Female	0.096154	0.008197	2
4	Female	0.250000	0.016393	2

In [11]:



```
km=KMeans()
```

In [12]:



```
y_predicted=km.fit_predict(df[["Age", "Income($)"]])
y_predicted
```

C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

```
warnings.warn(
```

Out[12]:

```
array([0, 0, 0, 0, 7, 0, 7, 0, 5, 7, 5, 7, 2, 0, 7, 0, 7, 0, 2, 7, 7, 0,
       2, 7, 2, 7, 2, 7, 7, 0, 5, 0, 2, 0, 2, 0, 2, 7, 7, 0, 5, 0, 2, 7,
       2, 0, 2, 7, 7, 7, 2, 7, 7, 5, 2, 2, 2, 5, 3, 2, 5, 3, 5, 2, 5, 3,
       2, 5, 3, 7, 5, 2, 5, 5, 5, 3, 2, 2, 3, 2, 5, 4, 5, 2, 3, 2, 1, 3,
       4, 1, 5, 3, 1, 4, 4, 3, 1, 3, 1, 3, 3, 1, 5, 3, 1, 3, 5, 1, 5, 5,
       5, 3, 4, 3, 3, 3, 5, 1, 1, 1, 3, 4, 4, 4, 3, 4, 1, 4, 1, 4, 1, 4,
       3, 4, 3, 4, 1, 4, 3, 4, 1, 4, 4, 4, 3, 4, 1, 4, 4, 4, 1, 4, 1, 4,
       1, 4, 4, 4, 4, 4, 1, 4, 3, 4, 1, 4, 4, 4, 4, 4, 4, 4, 4, 1, 4,
       1, 4, 1, 4, 6, 6, 1, 6, 6, 6, 1, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6,
       6, 6])
```

In [13]:



```
df["New Cluster"]=y_predicted
df.head()
```

Out[13]:


	Gender	Age	Income(\$)	cluster	New Cluster
0	Male	0.019231	0.000000	2	0
1	Male	0.057692	0.000000	2	0
2	Female	0.038462	0.008197	2	0
3	Female	0.096154	0.008197	2	0
4	Female	0.250000	0.016393	2	7

In [14]:

```
df1=df[df["New Cluster"]==0]
df2=df[df["New Cluster"]==1]
df3=df[df["New Cluster"]==2]
plt.scatter(df1["Age"],df1["Income($)"],color="red")
plt.scatter(df2["Age"],df2["Income($)"],color="green")
plt.scatter(df3["Age"],df3["Income($)"],color="blue")
plt.xlabel("Age")
plt.ylabel("Income($)")
```

Out[14]:

Text(0, 0.5, 'Income(\$)')



In [15]:

```
km.cluster_centers_
```

Out[15]:

```
array([[0.07239819, 0.08003857],
       [0.62037037, 0.47996357],
       [0.58974359, 0.20969945],
       [0.07322485, 0.38272383],
       [0.30944056, 0.50428465],
       [0.89799331, 0.28011404],
       [0.32905983, 0.78551913],
       [0.28388278, 0.1245121 ]])
```

In [16]:

```
km.cluster_centers_
```

Out[16]:


```
array([[0.07239819, 0.08003857],  
       [0.62037037, 0.47996357],  
       [0.58974359, 0.20969945],  
       [0.07322485, 0.38272383],  
       [0.30944056, 0.50428465],  
       [0.89799331, 0.28011404],  
       [0.32905983, 0.78551913],  
       [0.28388278, 0.1245121 ]])
```

In [17]:

```
df1=df[df["New Cluster"]==0]  
df2=df[df["New Cluster"]==1]  
df3=df[df["New Cluster"]==2]  
plt.scatter(df1["Age"],df1["Income($)"],color="red")  
plt.scatter(df2["Age"],df2["Income($)"],color="green")  
plt.scatter(df3["Age"],df3["Income($)"],color="blue")  
plt.scatter(km.cluster_centers_[0],km.cluster_centers_[1],color="orange",marker="+")  
plt.xlabel("Age")  
plt.ylabel("Income($)")
```

Out[17]:

```
Text(0, 0.5, 'Income($)')
```



In [18]:



```
k_rng=range(1,10)  
sse=[]
```

In [21]:



```
for k in k_rng:
    km=KMeans(n_clusters=k)
    km.fit(df[["Age", "Income($)"]])
    sse.append(km.inertia_)
#km.inertia_ will give you the value of sum of square error
print(sse)
plt.plot(k_rng, sse)
plt.xlabel("K")
plt.ylabel("Sum of Squared Error")
```

```
C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
```

```
[23.583906150363607, 13.028938428018286, 7.492113413237458, 6.058372453353154, 4.742998936528983, 3.8580680007628616, 3.0547174363693586, 2.645864014245702, 2.3135720353543285]
```

Out[21]:

Text(0, 0.5, 'Sum of Squared Error')



In []: